
The R Book

Second Edition

Michael J. Crawley

Imperial College London at Silwood Park, UK

<http://www.bio.ic.ac.uk/research/mjcraw/therbook/index.htm>



WILEY

A John Wiley & Sons, Ltd., Publication

This edition first published 2013
© 2013 John Wiley & Sons, Ltd

Registered office

John Wiley & Sons Ltd, The Atrium, Southern Gate, Chichester, West Sussex PO19 8SQ, United Kingdom

For details of our global editorial offices, for customer services and for information about how to apply for permission to reuse the copyright material in this book please see our website at www.wiley.com.

The right of the author to be identified as the author of this work has been asserted in accordance with the Copyright, Designs and Patents Act 1988.

All rights reserved. No part of this publication may be reproduced, stored in a retrieval system, or transmitted, in any form or by any means, electronic, mechanical, photocopying, recording or otherwise, except as permitted by the UK Copyright, Designs and Patents Act 1988, without the prior permission of the publisher.

Wiley also publishes its books in a variety of electronic formats. Some content that appears in print may not be available in electronic books.

Designations used by companies to distinguish their products are often claimed as trademarks. All brand names and product names used in this book are trade names, service marks, trademarks or registered trademarks of their respective owners. The publisher is not associated with any product or vendor mentioned in this book. This publication is designed to provide accurate and authoritative information in regard to the subject matter covered. It is sold on the understanding that the publisher is not engaged in rendering professional services. If professional advice or other expert assistance is required, the services of a competent professional should be sought.

Library of Congress Cataloging-in-Publication Data

Crawley, Michael J.

The R book / Michael J. Crawley. – 2e.

pages cm

Includes bibliographical references and index.

ISBN 978-0-470-97392-9 (hardback)

1. R (Computer program language) 2. Mathematical statistics—Data processing. I. Title.

QA276.45.R3C73 2013

519.50285'5133—dc23

2012027339

A catalogue record for this book is available from the British Library.

ISBN: 978-0-470-97392-9

Set in 10/12pt Times by Aptara Inc., New Delhi, India.

Mixed-Effects Models

Up to this point, we have treated all categorical explanatory variables as if they were the same. This is certainly what R.A. Fisher had in mind when he invented the analysis of variance in the 1920s and 1930s. It was Eisenhart (1947) who realized that there were actually two fundamentally different sorts of categorical explanatory variables: he called these **fixed effects** and **random effects**. It will take a good deal of practice before you are confident in deciding whether a particular categorical explanatory variable should be treated as a fixed effect or a random effect, but in essence:

- fixed effects influence only the **mean** of y ;
- random effects influence only the **variance** of y .

Fixed effects are unknown constants to be estimated from the data. Random effects govern the variance–covariance structure of the response variable (see p. 519). Nesting (or hierarchical structure) of random effects is a classic source of pseudoreplication, so it is important that you are able to recognize it and hence not fall into its trap. Random effects that come from the same group will be correlated, and this contravenes one of the fundamental assumptions of standard statistical models: *independence of errors*. Random effects occur in two contrasting kinds of circumstances:

- observational studies with hierarchical structure;
- designed experiments with different spatial or temporal scales.

Fixed effects have informative factor levels, while random effects often have uninformative factor levels. The distinction is best seen by an example. In most mammal species the categorical variable sex has two levels: male and female. For any individual that you find, the knowledge that it is, say, female conveys a great deal of information about the individual, and this information draws on experience gleaned from many other individuals that were female. A female will have a whole set of attributes (associated with her being female) no matter what population that individual was drawn from. Take a different categorical variable like genotype. If we have two genotypes in a population we might label them A and B. If we take two more genotypes from a *different* population we might label them A and B as well. In a case like this, the label A does not convey any information at all about the genotype, other than that it is probably different from genotype B. In the case of sex, the factor level (male or female) is informative: sex is a fixed effect. In the case of genotype, the factor level (A or B) is uninformative: genotype is a random effect.

Random effects have factor levels that are drawn from a large (potentially very large) population in which the individuals differ in many ways, but we do not know exactly how or why they differ. To get a feel for the difference between fixed effects and random effects here are some more examples:

Fixed effects	Random effects
Drug administered or not	Genotype
Insecticide sprayed or not	Brood
Nutrient added or not	Block within a field
One country versus another	Split plot within a plot
Male or female	History of development
Upland or lowland	Household
Wet versus dry	Individuals with repeated measures
Light versus shade	Family
One age versus another	Parent

The important point is that because the random effects come from a large population, there is not much point in concentrating on estimating means of our small subset of factor levels, and no point at all in comparing individual pairs of means for different factor levels. Much better to recognize them for what they are, random samples from a much larger population, and to concentrate on their variance. This is the *added* variation caused by differences between the levels of the random effects.

Variance components analysis is all about estimating the size of this variance, and working out its percentage contribution to the overall variation. There are five fundamental assumptions of linear mixed-effects models:

- Within-group errors are independent with mean zero and variance σ^2 .
- Within-group errors are independent of the random effects.
- The random effects are normally distributed with mean zero and covariance matrix Ψ .
- The random effects are independent in different groups.
- The covariance matrix does not depend on the group.

The validity of these assumptions needs to be tested by employing a series of plotting methods involving the residuals, the fitted values and the predicted random effects. The tricks with mixed-effects models are:

- learning which variables are random effects;
- specifying the fixed and random effects in the model formula;
- getting the nesting structure of the random effects right;
- remembering to get `library(lme4)` or `library(nlme)` at the outset.

The issues fall into two broad categories: questions about experimental design and the management of experimental error (e.g. where does most of the variation occur, and where would increased replication be most profitable?); and questions about hierarchical structure, and the relative magnitude of variation at different levels within the hierarchy (e.g. studies on the genetics of individuals within families, families within parishes, and parishes with counties, to discover the relative importance of genetic and phenotypic variation).

Most ANOVA models are based on the assumption that there is a single error term. But in hierarchical studies and nested experiments, where the data are gathered at two or more different spatial scales, there is *a different error variance for each different spatial scale*. There are two reasonably clear-cut sets of circumstances where your first choice would be to use a linear mixed-effects model: you want to do variance components analysis because all your explanatory variables are categorical random effects and you do not have any fixed effects; or you do have fixed effects, but you also have pseudoreplication of one sort or another (e.g. temporal pseudoreplication resulting from repeated measurements on the same individuals; see p. 699). To test whether one should use a model with mixed effects or just a plain old linear model, Douglas Bates wrote in the R help archive: ‘I would recommend the likelihood ratio test against a linear model fit by `lm`. The *p*-value returned from this test will be conservative because you are testing on the boundary of the parameter space.’

19.1 Replication and pseudoreplication

To qualify as replicates, measurements must have the following properties:

- They must be independent.
- They must not form part of a time series (data collected from the same place on successive occasions are not independent).
- They must not be grouped together in one place (aggregating the replicates means that they are not spatially independent).
- They must be of an appropriate spatial scale;
- Ideally, one replicate from each treatment ought to be grouped together into a block, and the whole experiment repeated in many different blocks.
- Repeated measures (e.g. from the same individual or the same spatial location) are not replicates (this is probably the commonest cause of pseudoreplication in statistical work).

Pseudoreplication occurs when you analyse the data as if you had more degrees of freedom than you really have. There are two kinds of pseudoreplication:

- temporal pseudoreplication, involving repeated measurements from the same individual;
- spatial pseudoreplication, involving several measurements taken from the same vicinity.

Pseudoreplication is a problem because one of the most important assumptions of standard statistical analysis is *independence of errors*. Repeated measures through time on the same individual will have non-independent errors because peculiarities of the individual will be reflected in all of the measurements made on it (the repeated measures will be temporally correlated with one another). Samples taken from the same vicinity will have non-independent errors because peculiarities of the location will be common to all the samples (e.g. yields will all be high in a good patch and all be low in a bad patch).

Pseudoreplication is generally quite easy to spot. The question to ask is this. How many degrees of freedom for error does the experiment really have? If a field experiment appears to have lots of degrees of freedom, it is probably pseudoreplicated. Take an example from pest control of insects on plants. There are 20 plots, 10 sprayed and 10 unsprayed. Within each plot there are 50 plants. Each plant is measured five times during

the growing season. Now this experiment generates $20 \times 50 \times 5 = 5000$ numbers. There are two spraying treatments, so there must be 1 degree of freedom for spraying and 4998 degrees of freedom for error. Or must there? Count up the replicates in this experiment. Repeated measurements on the same plants (the five sampling occasions) are certainly not replicates. The 50 individual plants within each quadrat are not replicates either. The reason for this is that conditions within each quadrat are quite likely to be unique, and so all 50 plants will experience more or less the same unique set of conditions, irrespective of the spraying treatment they receive. In fact, there are 10 replicates in this experiment. There are 10 sprayed plots and 10 unsprayed plots, and each plot will yield only one independent datum for the response variable (the mean proportion of leaf area consumed by insects, for example). Thus, there are 9 degrees of freedom within each treatment, and $2 \times 9 = 18$ degrees of freedom for error in the experiment as a whole. It is not difficult to find examples of pseudoreplication on this scale in the literature (Hurlbert 1984). The problem is that it leads to the reporting of masses of spuriously significant results (with 4998 degrees of freedom for error, it is almost impossible *not* to have significant differences). The first skill to be acquired by the budding experimenter is the ability to plan an experiment that is properly replicated. There are various things that you can do when your data are pseudoreplicated:

- Average away the pseudoreplication and carry out your statistical analysis on the means.
- Carry out separate analyses for each time period.
- Use proper time series analysis or mixed-effects models.

19.2 The `lme` and `lmer` functions

Most of the examples in this chapter use the linear mixed model formula `lme`. This is to provide compatibility with the excellent book by Pinheiro and Bates (2000) on *Mixed-Effects Models in S and S-PLUS*. More recently, however, Douglas Bates has released the generalized mixed model function `lmer` as part of the `lme4` package, and you may prefer to use this in your own work, especially for nested count data or proportion data. To begin with, I provide a simple comparison of the basic syntax of the two functions.

19.2.1 `lme`

Specifying the fixed and random effects in the model formula is done with two formulae. Suppose that there are no fixed effects, so that all of the categorical variables are random effects. Then the fixed effect simply estimates the intercept (parameter 1):

```
fixed = y ~ 1
```

The fixed effect (a compulsory part of the `lme` structure) is just the overall mean value of the response variable `y ~ 1`. The `fixed =` part of the formula is optional if you put this object first. The random effects show the identities of the random variables and their relative locations in the hierarchy. The three random effects (`a`, `b`, and `c`) are specified like this:

```
random = ~ 1 | a/b/c
```

and in this case the phrase `random =` is *not* optional. An important detail to notice is that the name of the response variable (`y`) is not repeated in the random-effects formula: there is a blank space to the left of the tilde `~`. In most mixed-effects models we assume that the random effects have a mean of zero and that we are interested in quantifying variation in the intercept caused by differences between the factor levels of the

random effects. After the intercept comes the vertical bar `|` which is read as ‘given the following spatial arrangement of the random variables’. In this example there are three random effects with ‘c nested within b which in turn is nested within a’. The factors are separated by forward slash characters, and the variables are listed from left to right in declining order of spatial (or temporal) scale. This will only become clear with practice, but it is a simple idea. The formulae are put together like this:

```
lme(fixed = y ~ 1, random = ~ 1 | a/b/c)
```

19.2.2 lmer

There is just one formula in `lmer`, not separate formulae for the fixed and random effects. The fixed effects are specified first, to the right of the tilde, in the normal way. Next comes a plus sign, then one or more random terms enclosed in parentheses (in this example there is just one random term, but we might want separate random terms for the intercept and for the slopes, for instance). R can identify the random terms because they must contain a ‘given’ symbol `|`, to the right of which are listed the random effects in the usual way, from largest to smallest scale, left to right. So the `lmer` formula for this example is:

```
lmer(y ~ 1 + ( 1 | a/b/c ))
```

19.3 Best linear unbiased predictors

In `aoov`, the effect size for treatment i is defined as $\bar{y}_i - \mu$, where μ is the overall mean. In mixed-effects models, however, correlation between the pseudoreplicates within a group causes what is called **shrinkage**. The best linear unbiased predictors (BLUPs, denoted by a_i) are smaller than the effect sizes ($\bar{y}_i - \mu$), and are given by

$$a_i = (\bar{y}_i - \mu) \left(\frac{\sigma_a^2}{\sigma_a^2 + \sigma^2/n} \right),$$

where σ^2 is the residual variance and σ_a^2 is the between-group variance which introduces the correlation between the pseudoreplicates within each group. Thus, the parameter estimate a_i is ‘shrunk’ compared to the fixed effect size ($\bar{y}_i - \mu$). When σ_a^2 is estimated to be large compared with the estimate of σ^2/n (i.e. when most of the variation is between classes and there is little variation within classes), the fixed effects and the BLUP are similar. On the other hand, when σ_a^2 is estimated to be small compared with the estimate of σ^2/n , then the fixed effects and the BLUP can be very different.

19.4 Designed experiments with different spatial scales: Split plots

The important distinction in models with categorical explanatory variables is between cases where the data come from a designed experiment, in which treatments were allocated to locations or subjects at random, and cases where the data come from an observational study in which the categorical variables are associated with an observation before the study. Here, we call the first case split-plot experiments and the second case hierarchical designs. The point is that their dataframes look identical, so it is easy to analyse one case wrongly as if it were the other. You need to be able to distinguish between fixed effects and random effects in both cases.

Here is the linear model for a split-plot experiment analysed in Chapter 11 by `aov` (see p. 519):

```
yields <- read.table("c:\\temp\\splityield.txt",header=T)
attach(yields)
names(yields)

[1] "yield" "block" "irrigation" "density" "fertilizer"

library(nlme)
```

The fixed-effects part of the model is specified in just the same way as in a straightforward factorial experiment: `yield ~ irrigation*density*fertilizer`. The random-effects part of the model says that we want the random variation to enter via effects on the intercept as `random=~1`. Finally, we define the spatial structure of the random effects after the ‘given’ symbol `|` as: `block/irrigation/density` reflecting the progressively smaller plot sizes. There is no need to specify the smallest spatial scale (fertilizer plots in this example).

```
model <- lme(yield ~ irrigation*density*fertilizer,random= ~
1|block/irrigation/density)
summary(model)
```

Linear mixed-effects model fit by REML

```
Data: NULL
      AIC      BIC    logLik
481.6212 525.3789 -218.8106
```

Random effects:

```
Formula: ~1 | block
(Intercept)
```

StdDev: 0.000660972

```
Formula: ~1 | irrigation %in% block
(Intercept)
```

StdDev: 1.982463

```
Formula: ~1 | density %in% irrigation %in% block
(Intercept) Residual
```

StdDev: 6.975553 9.292805

Fixed effects: `yield ~ irrigation * density * fertilizer`

	Value	Std.Error	DF	t-value	p-value
(Intercept)	80.50	5.893741	36	13.658558	0.0000
irrigationirrigated	31.75	8.335008	3	3.809234	0.0318
densitylow	5.50	8.216281	12	0.669403	0.5159
densitymedium	14.75	8.216281	12	1.795216	0.0978
fertilizerNP	5.50	6.571005	36	0.837010	0.4081
fertilizerP	4.50	6.571005	36	0.684827	0.4978
irrigationirrigated:densitylow	-39.00	11.619577	12	-3.356405	0.0057
irrigationirrigated:densitymedium	-22.25	11.619577	12	-1.914872	0.0796
irrigationirrigated:fertilizerNP	13.00	9.292805	36	1.398932	0.1704
irrigationirrigated:fertilizerP	5.50	9.292805	36	0.591856	0.5576
densitylow:fertilizerNP	3.25	9.292805	36	0.349733	0.7286
densitymedium:fertilizerNP	-6.75	9.292805	36	-0.726368	0.4723
densitylow:fertilizerP	-5.25	9.292805	36	-0.564953	0.5756
densitymedium:fertilizerP	-5.50	9.292805	36	-0.591856	0.5576
irrigationirrigated:densitylow:fertilizerNP	7.75	13.142011	36	0.589712	0.5591
irrigationirrigated:densitymedium:fertilizerNP	3.75	13.142011	36	0.285344	0.7770
irrigationirrigated:densitylow:fertilizerP	20.00	13.142011	36	1.521837	0.1368
irrigationirrigated:densitymedium:fertilizerP	4.00	13.142011	36	0.304367	0.7626

This output suggests that the only significant effects are the main effect of `irrigation` ($p = 0.0318$) and the `irrigation` by `density` interaction ($p = 0.0057$). The three-way interaction is not significant so we remove it, fitting all terms up to two-way interactions:

```
model <- lme(yield~(irrigation+density+fertilizer)^2,
             random=~1|block/irrigation/density)
summary(model)
```

Linear mixed-effects model fit by REML
Data: NULL

	AIC	BIC	logLik
	503.1256	540.2136	-233.5628

Random effects:

Formula: ~1 | block
(Intercept)
StdDev: 0.0005634512

Formula: ~1 | irrigation %in% block
(Intercept)
StdDev: 1.982562

Formula: ~1 | density %in% irrigation %in% block
(Intercept) Residual
StdDev: 7.041303 9.142696

Fixed effects: yield ~ (irrigation + density + fertilizer)^2

	Value	Std.Error	DF	t-value	p-value
(Intercept)	82.47222	5.443438	40	15.150760	0.0000
irrigationirrigated	27.80556	7.069256	3	3.933307	0.0293
densitylow	0.87500	7.256234	12	0.120586	0.9060
densitymedium	13.45833	7.256234	12	1.854727	0.0884
fertilizerNP	3.58333	5.278538	40	0.678850	0.5011
fertilizerP	0.50000	5.278538	40	0.094723	0.9250
irrigationirrigated:densitylow	-29.75000	8.800165	12	-3.380618	0.0055
irrigationirrigated:densitymedium	-19.66667	8.800165	12	-2.234807	0.0452
irrigationirrigated:fertilizerNP	16.83333	5.278538	40	3.189014	0.0028
irrigationirrigated:fertilizerP	13.50000	5.278538	40	2.557526	0.0144
densitylow:fertilizerNP	7.12500	6.464862	40	1.102112	0.2770
densitymedium:fertilizerNP	-4.87500	6.464862	40	-0.754076	0.4552
densitylow:fertilizerP	4.75000	6.464862	40	0.734741	0.4668
densitymedium:fertilizerP	-3.50000	6.464862	40	-0.541388	0.5912

The fertilizer by density interaction is not significant, so we remove it:

```
model <- lme(yield~irrigation*density+irrigation*fertilizer,
             random=~1|block/irrigation/density)
summary(model)
```

Linear mixed-effects model fit by REML
Data: NULL

	AIC	BIC	logLik
	519.9035	549.6834	-245.9517

Random effects:

Formula: ~1 | block

```

      (Intercept)
StdDev: 0.0005566885

Formula: ~1 | irrigation %in% block
      (Intercept)
StdDev:      1.982614

Formula: ~1 | density %in% irrigation %in% block
      (Intercept) Residual
StdDev:      7.057132 9.105995

Fixed effects: yield ~ irrigation * density + irrigation * fertilizer

              Value Std.Error DF   t-value p-value
(Intercept)    82.08333   4.994999 44  16.433103  0.0000
irrigationirrigated 27.80556   7.063995  3   3.936236  0.0292
densitylow         4.83333   6.222653 12   0.776732  0.4524
densitymedium     10.66667   6.222653 12   1.714167  0.1122
fertilizerNP       4.33333   3.717507 44   1.165656  0.2500
fertilizerP        0.91667   3.717507 44   0.246581  0.8064
irrigationirrigated:densitylow -29.75000   8.800161 12  -3.380620  0.0055
irrigationirrigated:densitymedium -19.66667   8.800161 12  -2.234808  0.0452
irrigationirrigated:fertilizerNP 16.83333   5.257349 44   3.201867  0.0025
irrigationirrigated:fertilizerP 13.50000   5.257349 44   2.567834  0.0137

```

Both the `irrigation` by `fertilizer` and `irrigation` by `density` interactions are now highly significant. The apparently non-significant main effect of `density` is spurious because `density` appears in a significant interaction with `irrigation`. The moral is that you must do the model simplification to get the appropriate *p* values.

Remember, too, that if you want to use `anova` to compare mixed models with different fixed-effects structures, then you must use maximum likelihood (`method = "ML"` in `lme` but `REML = FALSE` in `lmer`) rather than the default restricted maximum likelihood (REML). Here is the analysis again, but this time using `anova` to compare models with progressively simplified fixed effects:

```

model.lme <- lme(yield~irrigation*density*fertilizer,
                 random=~ 1| block/irrigation/density,method="ML")
model.lme.2 <- update(model.lme,~. - irrigation:density:fertilizer)
anova(model.lme,model.lme.2)

      Model df      AIC      BIC    logLik   Test  L.Ratio p-value
model.lme      1 22 573.5108 623.5974 -264.7554
model.lme.2     2 18 569.0046 609.9845 -266.5023 1 vs 2 3.493788 0.4788

model.lme.3 <- update(model.lme.2,~. - density:fertilizer)
anova(model.lme.3,model.lme.2)

      Model df      AIC      BIC    logLik   Test  L.Ratio p-value
model.lme.3     1 14 565.1933 597.0667 -268.5967
model.lme.2     2 18 569.0046 609.9845 -266.5023 1 vs 2 4.188774 0.3811

model.lme.4 <- update(model.lme.3,~. - irrigation:fertilizer)
anova(model.lme.3,model.lme.4)

```

	Model	df	AIC	BIC	logLik	Test	L.Ratio	p-value
model.lme.3	1	14	565.1933	597.0667	-268.5967			
model.lme.4	2	12	572.3373	599.6573	-274.1687	1 vs 2	11.14397	0.0038

```
model.lme.5 <- update(model.lme.2, ~. - irrigation:density)
anova(model.lme.5, model.lme.2)
```

	Model	df	AIC	BIC	logLik	Test	L.Ratio	p-value
model.lme.5	1	16	576.7134	613.1400	-272.3567			
model.lme.2	2	18	569.0046	609.9845	-266.5023	1 vs 2	11.70883	0.0029

The `irrigation` by `fertilizer` interaction is more significant ($p = 0.0038$ compared to $p = 0.0081$) under this mixed-effects model than it was in the linear model earlier, as is the `irrigation` by `density` interaction ($p = 0.0029$ compared to $p = 0.01633$). You need to do the model simplification in `lme` to uncover the significance of the main effect and interaction terms, but it is worth it, because the `lme` analysis can be more powerful. The minimal adequate model under the `lme` is:

```
summary(model.lme.3)
```

Linear mixed-effects model fit by maximum likelihood

Data: NULL

	AIC	BIC	logLik
	565.1933	597.0667	-268.5967

Random effects:

Formula: ~1 | block
(Intercept)

StdDev: 0.0005260885

Formula: ~1 | irrigation %in% block
(Intercept)

StdDev: 1.716888

Formula: ~1 | density %in% irrigation %in% block
(Intercept) Residual

StdDev: 5.722413 8.718327

Fixed effects: yield ~ irrigation + density + fertilizer

+ irrigation:density + irrigation:fertilizer

	Value	Std.Error	DF	t-value	p-value
(Intercept)	82.08333	4.756285	44	17.257867	0.0000
irrigationirrigated	27.80556	6.726403	3	4.133793	0.0257
densitylow	4.83333	5.807347	12	0.832279	0.4215
densitymedium	10.66667	5.807347	12	1.836754	0.0911
fertilizerNP	4.33333	3.835552	44	1.129781	0.2647
fertilizerP	0.91667	3.835552	44	0.238992	0.8122
irrigationirrigated:densitylow	-29.75000	8.212829	12	-3.622382	0.0035
irrigationirrigated:densitymedium	-19.66667	8.212829	12	-2.394628	0.0338
irrigationirrigated:fertilizerNP	16.83333	5.424290	44	3.103325	0.0033
irrigationirrigated:fertilizerP	13.50000	5.424290	44	2.488805	0.0167

You should pay special attention to the degrees of freedom column. Note that the degrees of freedom are not pseudoreplicated: there are only 3 d.f. for testing the `irrigation` main effect; 12 d.f. for testing the `irrigation` by `density` interaction and 44 d.f. for `irrigation` by `fertilizer` (this is $36 + 4 + 4 = 44$ after model simplification). Also, remember that you must do your model simplification using maximum likelihood (`method = "ML"`) because you cannot use `anova` to compare models with different fixed-effect structures using REML.

Model-checking plots show that the residuals are well behaved:


```
plot(model.lme.3)
```

the response variable is a reasonably linear function of the fitted values:

```
plot(model.lme.3, yield ~ fitted(.))
```

and the errors are reasonably close to normally distributed in all four blocks:

```
qqnorm(model.lme.3, ~ resid(.) | block)
```



When, as here, the experiment is balanced and there are no missing values, then it is much simpler to interpret the `aov` using an `Error` term to describe the structure of the spatial pseudoreplication (p. 526), not least because it produces separate ANOVA tables for each of the spatial scales, which makes it very easy to check that there is no pseudoreplication. Without balance, however, you will need to use `lme` or `lmer` and to use model simplification to estimate the *p* values of the significant interaction terms.

If you do this example using `lmer`, you will want to switch off the matrix of correlations for the fixed effects. You do this with the `print(model, cor=F)` option (rather than `summary`):

```
library(lme4)
```

```
b <- block
```

```
bi <- block:irrigation
```

```

bid <- block:irrigation:density
modell <-
lmer(yield~irrigation*density*fertilizer+(1|b)+(1|bi)+(1|bid),REML=FALSE)
print(modell,cor=F)

```

Linear mixed model fit by maximum likelihood

Formula: yield ~ irrigation * density * fertilizer + (1|b) + (1|bi) + (1|bid)

	AIC	BIC	logLik	deviance	REMLdev
	573.5	623.6	-264.8	529.5	437.6

Random effects:

Groups	Name	Variance	Std.Dev.
bid	(Intercept)	3.6493e+01	6.0410e+00
bi	(Intercept)	2.9479e+00	1.7169e+00
b	(Intercept)	8.9145e-13	9.4417e-07
Residual		6.4767e+01	8.0478e+00

Number of obs: 72, groups: bid, 24; bi, 8; b, 4

Fixed effects:

	Estimate	Std. Error	t value
(Intercept)	80.500	5.104	15.772
irrigationirrigated	31.750	7.218	4.399
densitylow	5.500	7.115	0.773
densitymedium	14.750	7.115	2.073
fertilizerNP	5.500	5.691	0.966
fertilizerP	4.500	5.691	0.791
irrigationirrigated:densitylow	-39.000	10.063	-3.876
irrigationirrigated:densitymedium	-22.250	10.063	-2.211
irrigationirrigated:fertilizerNP	13.000	8.048	1.615
irrigationirrigated:fertilizerP	5.500	8.048	0.683
densitylow:fertilizerNP	3.250	8.048	0.404
densitymedium:fertilizerNP	-6.750	8.048	-0.839
densitylow:fertilizerP	-5.250	8.048	-0.652
densitymedium:fertilizerP	-5.500	8.048	-0.683
irrigationirrigated:densitylow:fertilizerNP	7.750	11.381	0.681
irrigationirrigated:densitymedium:fertilizerNP	3.750	11.381	0.329
irrigationirrigated:densitylow:fertilizerP	20.000	11.381	1.757
irrigationirrigated:densitymedium:fertilizerP	4.000	11.381	0.351

As before, it requires model simplification before the significant interactions become evident.

19.5 Hierarchical sampling and variance components analysis

Hierarchical data are often encountered in observational studies where information is collected at a range of different spatial scales. The principal aim is to discover the scale at which most of the variation is generated. This information would then allow a closer focus on mechanisms operating at this scale in subsequent more detailed studies. The following study involves a test with a mean score of 100 administered to children in four British towns. Each town was divided into districts by postcodes, and six districts were selected at random. Within districts, 10 streets were selected at random, and within streets, four households were

selected at random. Naturally, different households had different numbers of children (childless households were excluded from the study) and there was no control over sex ratio of children within household.

```
library(lme4)

data <- read.table("c:\\temp\\childfull.txt",header=T)
attach(data)
head(data)
```

	childID	child	house	street	district	town	response	gender
1	1	1	door1	1	A Leeds	83.88773	male	
2	1	1	door2	1	A Leeds	99.96294	male	
3	1	3	door3	1	A Leeds	87.20253	female	
4	2	3	door3	1	A Leeds	89.37665	male	
5	3	3	door3	1	A Leeds	92.01751	female	
6	1	5	door4	1	A Leeds	87.12672	female	

You can see that the factor levels are not unique: for instance, there is a street 1 in each district of Leeds (and of every other town). We use the colon operator to create unique factor levels for each random effect: district within town (*d*), street within district within town (*s*) and household within street within district within town (*h*). Each household has one or more children (maximum = 8, mostly in Leeds) but the sex ratio varies from house to house.

```
d <- town:district
s <- town:district:factor(street)
h <- town:district:factor(street):house
```

The mixed effects model has one fixed effect (gender) and four nested random effects:

```
model <- lmer(response~gender+(1|town)+(1|d)+(1|s)+(1|h))
summary(model)
```

```
Linear mixed model fit by REML
Formula: response ~ gender + (1 | town) + (1 | d) + (1 | s) + (1 | h)
      AIC      BIC logLik deviance REMLdev
19878 19920   -9932   19868   19864
Random effects:
Groups      Name          Variance Std.Dev.
h           (Intercept)    4.0817  2.0203
s           (Intercept)   15.6747  3.9591
d           (Intercept)  168.3451 12.9748
town        (Intercept)   36.9802  6.0811
Residual                    36.2405  6.0200

Number of obs: 2972, groups: h, 960; s, 240; d, 24; town, 4

Fixed effects:
              Estimate Std. Error t value
(Intercept)  97.8965     4.0424  24.218
gendermale    0.5368     0.2363   2.272

Correlation of Fixed Effects:
              (Intr)
gendermale -0.030
```

The fixed effect is significant ($t = 2.272$) but small (0.537) compared to the means for towns, which varied by as much as 20 (e.g. Coventry vs. Derby). Of the random effects, most of the variation in the response is between districts within towns (64%) and least is between households within streets (4%). You get the percentage variance components like this:

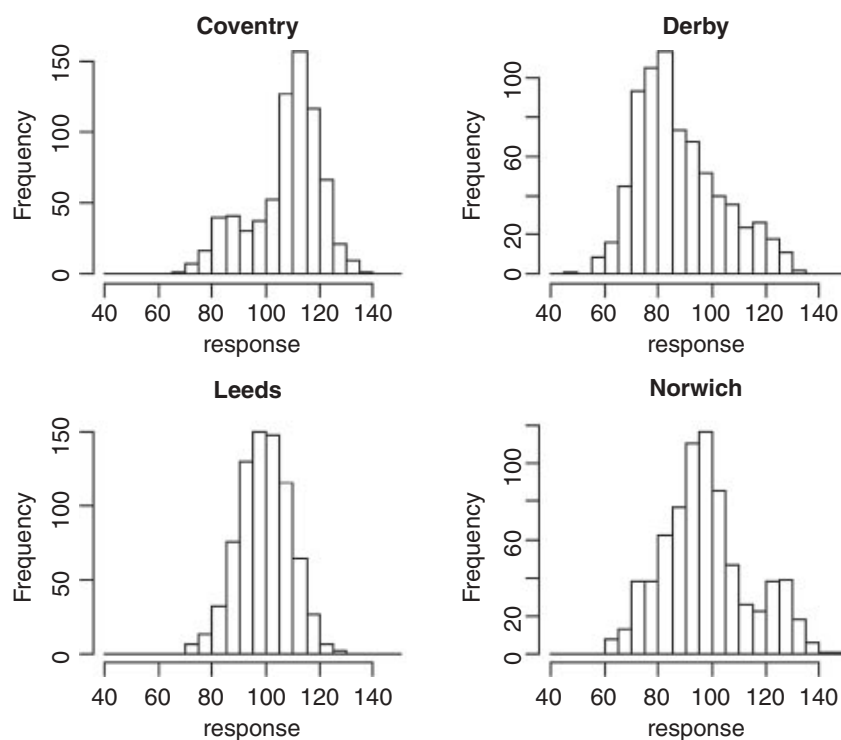
```
vc <- c(36.2405, 4.0817, 15.6747, 168.3451, 36.9802)
vc <- 100*c(36.2405, 4.0817, 15.6747, 168.3451, 36.9802) / sum(vc)
vc

[1] 13.868129  1.561942  5.998227 64.420512 14.151190
```

The key feature of these data, however, is the substantial variation between children within the same family (i.e. even after you have controlled for family and environment (town and district within town)).

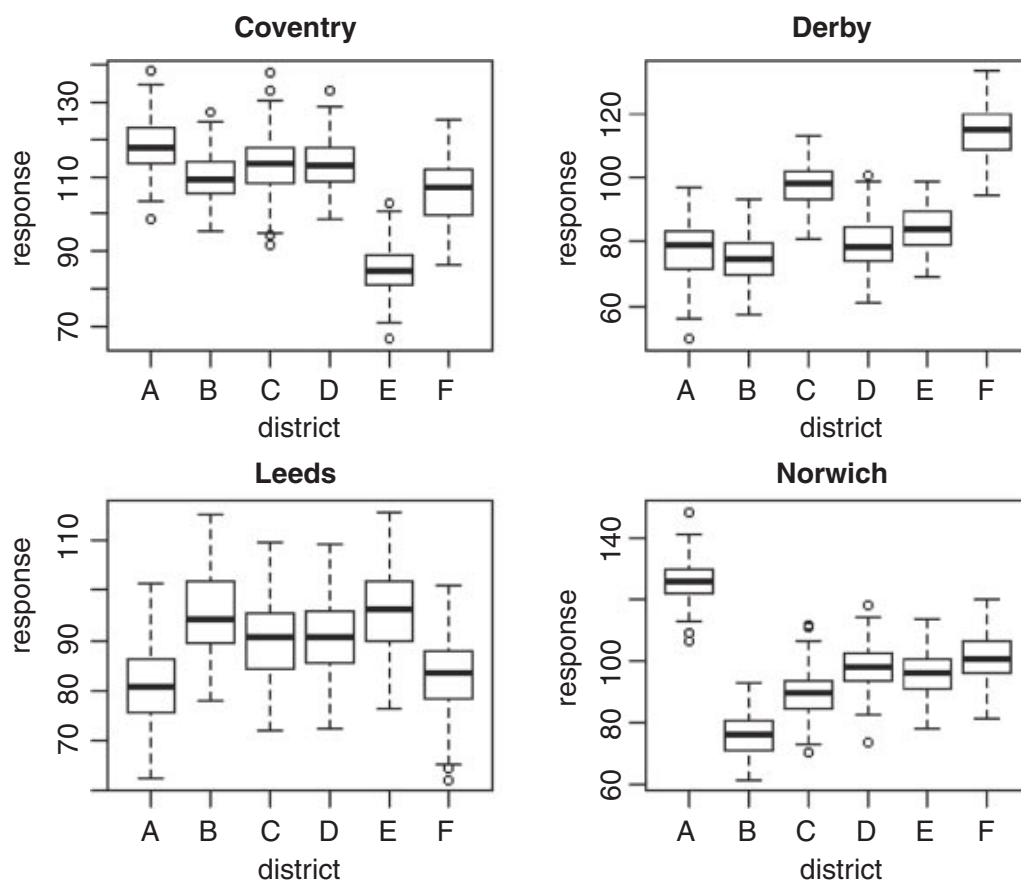
There are issues about the appropriate graphics in studies like this, simply because there are so many combinations of factor levels, and the nested factor levels typically make sense only in the context of their higher-level associations. Here is one way of showing the variation:

```
par(mfrow=c(2,2))
hist(response[town=="Coventry"], main="Coventry", breaks=seq(40, 150, 5),
     xlab="response")
hist(response[town=="Derby"], main="Derby", breaks=seq(40, 150, 5),
     xlab="response")
hist(response[town=="Leeds"], main="Leeds", breaks=seq(40, 150, 5),
     xlab="response")
hist(response[town=="Norwich"], main="Norwich", breaks=seq(40, 150, 5),
     xlab="response")
```




This is good at highlighting non-normality (e.g. the bimodal distribution of the response in Norwich, and the opposite skew in Coventry and Derby). Here is an alternative using box-and-whisker plots, with `subset` to chose the towns rather than subscripts:

```
plot(response~district,subset=(town=="Coventry"),main="Coventry")
plot(response~district,subset=(town=="Derby"),main="Derby")
plot(response~district,subset=(town=="Leeds"),main="Leeds")
plot(response~district,subset=(town=="Norwich"),main="Norwich")
```



This is particularly good for drawing attention to districts with strikingly different mean scores (e.g. high ones like district A in Norwich and low ones like district E in Coventry). For looking at interactions, `coplot` is often useful:

```
coplot(response~district|town)
```

In practice, you will want to try many different kinds of plots across all of the spatial scales.

19.6 Mixed-effects models with temporal pseudoreplication

A common cause of temporal pseudoreplication in growth experiments with fixed effects is when each individual is measured several times as it grows during the course of an experiment. The next example is as simple as possible: we have a single fixed effect (a two-level categorical variable, with fertilizer added or not) and six replicate plants in each treatment, with each plant measured on five occasions (after 2, 4, 6, 8 or 10 weeks of growth). The response variable is root length, measured non-destructively through a glass panel, which is opened to the light only when the root length measurements are being taken. The fixed-effect formula looks like this:

```
fixed = root~fertilizer
```

The random-effects formula needs to indicate that the week of measurement (a continuous random effect) represents pseudoreplication within each individual plant:

```
random = ~week|plant
```

Because we have a continuous random effect (weeks) we write `~week` in the random-effects formula rather than the `~1` that we used with categorical random effects (above). Here are the data:

```
results <- read.table("c:\\temp\\fertilizer.txt",header=T)
attach(results)
```

```
names(results)
```

```
[1] "root" "week" "plant" "fertilizer"
```

We begin with data inspection. For the kind of data involved in mixed-effects models there are some excellent built-in plotting functions (variously called panel plots, trellis plots, or lattice plots).

```
library(nlme)
```


```
library(lattice)
```

To use trellis plotting, we begin by turning our dataframe called `results` into a `groupedData` object (p. 957). To do this we specify the nesting structure of the random effects, and indicate the fixed effect by defining `fertilizer` as `outer` to this nesting (a fixed effect):

```
results <- groupedData(root~week|plant,outer = ~ fertilizer,results)
```

Because `results` is now a `groupedData` object, the plotting is fantastically simple:


```
plot(results)
```



Here you get separate time series plots for each of the individual plants (created, in this case, by joining the dots, which is the default option), with plant identities ranked from bottom left (ID5) to top right (ID7) on the basis of mean root length. In terms of understanding the fixed effects, it is often more informative to group together the six replicates within each treatment, and to have one panel for each of the treatment levels (i.e.

one for the fertilized plants and one for the controls in this case). This is very straightforward, using `outer` to indicate the grouping:

```
plot(results, outer=T)
```



You can see that by week 10 there is virtually no overlap between the two treatment groups. The largest control plant has about the same root length as the smallest fertilized plant (about 9 cm).

Now for the statistical modelling. Ignoring the pseudoreplication, we should have 1 d.f. for `fertilizer` and $2 \times (6 - 1) = 10$ d.f. for error.

```
model <- lme(root~fertilizer, random=~week|plant)
summary(model)
```

Linear mixed-effects model fit by REML

Data: NULL

	AIC	BIC	logLik
	171.0236	183.3863	-79.51181

Random effects:

Formula: ~week | plant

Structure: General positive-definite, Log-Cholesky parametrization

	StdDev	Corr
(Intercept)	2.8639832	(Intr)
week	0.9369412	-0.999
Residual	0.4966308	

Fixed effects: root ~ fertilizer

	Value	Std.Error	DF	t-value	p-value
(Intercept)	2.799710	0.1438367	48	19.464499	0e+00
fertilizercontrol	-1.039383	0.2034158	10	-5.109645	5e-04

Correlation:

	(Intr)
fertilizercontrol	-0.707

```
Standardized Within-Group Residuals:
      Min       Q1       Med       Q3       Max
-1.9928118 -0.6586834 -0.1004301  0.6949714  2.0225381

Number of Observations: 60
Number of Groups: 12
```

The output looks dauntingly complex, but once you learn your way around it, the essential information is relatively easy to extract. The mean reduction in root size associated with the unfertilized controls is $-1.039\,383$ and this has a standard error of $0.203\,415\,8$ based on the correct 10 residual d.f. (six replicates per factor level). Can you see why the intercept has 48 d.f.? (Hint: ask yourself how many graphs have been fitted to the data.)

Here is a simple one-way ANOVA for the non-pseudoreplicated data taken from the end of the experiment in week 10:

```
model2 <- aov(root~fertilizer,subset=(week==10))
summary(model2)
```

```
          Df Sum Sq Mean Sq F value Pr(>F)
fertilizer  1  4.941    4.941    11.49 0.0069 **
Residuals 10  4.302    0.430
```

```
summary.lm(model2)
```

Call:

```
aov(formula = root ~ fertilizer, subset = (week == 10))
```

Residuals:

```
      Min       1Q   Median       3Q      Max
-0.8167 -0.3667 -0.1333  0.4042  1.4833
```

Coefficients:

```
              Estimate Std. Error t value Pr(>|t|)
(Intercept)      9.6167    0.2678   35.915 6.65e-12 ***
fertilizercontrol -1.2833    0.3787   -3.389  0.0069 **
```

Residual standard error: 0.6559 on 10 degrees of freedom

Multiple R-squared: 0.5346, Adjusted R-squared: 0.488

F-statistic: 11.49 on 1 and 10 DF, p-value: 0.006897

We can compare this with the output from the `lme`. The effect size in the `lme` is slightly smaller ($-1.039\,393$ compared to -1.2833) but the standard error is appreciably lower ($0.203\,415\,8$ compared to 0.3787), so the significance of the result is higher in the `lme` than in the `aov`. You get increased statistical power as a result of going to the trouble of fitting the mixed-effects model. And, crucially, you do not need to make potentially arbitrary judgements about which time period to select for the non-pseudoreplicated analysis. You use all of the data in the model, and you specify its structure appropriately so that the hypotheses are tested with the correct degrees of freedom (10 in this case, not 48).

The reason why the effect sizes are different in the `lm` and `lme` models is that linear models use maximum likelihood estimates of the parameters based on arithmetic means. The linear mixed models, however, use the wonderfully named BLUPs (see p. 685).

19.7 Time series analysis in mixed-effects models

It is common to have repeated measures on subjects in observational studies, where we would expect that the observation on an individual at time $t + 1$ would be quite strongly correlated with the observation on the same individual at time t . This contravenes one of the central assumptions of linear models (p. 503), that the within-group errors are independent. However, we often observe significant serial correlation in data such as these.

The following example comes from Pinheiro and Bates (2000) and forms part of the `nlme` library. The data refer to the numbers of ovaries observed in repeated measures on 11 mares (their oestrus cycles have been scaled such that ovulation occurred at time 0 and at time 1). The issue is how best to model the correlation structure of the data. We know from previous work that the fixed effect can be modelled as a three-parameter sine–cosine function of time x :

$$y = a + b \sin(2\pi x) + d \cos(2\pi x) + \varepsilon_{ij},$$


and we want to assess different structures for modelling the within-class correlation.

The dataframe is of class `groupedData` which makes the plotting and error checking much simpler.

```
data(Ovary)
attach(Ovary)
names(Ovary)

[1] "Mare" "Time" "follicles"

plot(Ovary)
```



The panel plot has ranked the horses from bottom left to top right on the basis of their mean number of ovules (mare 4 with the lowest number, mare 8 with the highest). Some animals show stronger cyclic behaviour than others.

We begin by fitting a mixed-effects model making no allowance for the correlation structure, and investigate the degree of autocorrelation that is exhibited by the residuals (recall that the assumption of the model is that there is no correlation):

```
model <- lme(follicles~sin(2*pi*Time)+cos(2*pi*Time),
             data=Ovary,random=~ 1| Mare)
summary(model)
```

Linear mixed-effects model fit by REML

Data: Ovary			
AIC	BIC	logLik	
1669.36	1687.962	-829.6802	

Random effects:

Formula: ~1 Mare	
(Intercept)	Residual
StdDev: 3.041344	3.400466

Fixed effects: follicles ~ sin(2 * pi * Time) + cos(2 * pi * Time)

	Value	Std.Error	DF	t-value	p-value
(Intercept)	12.182244	0.9390009	295	12.973623	0.0000
sin(2 * pi * Time)	-3.339612	0.2894013	295	-11.539727	0.0000
cos(2 * pi * Time)	-0.862422	0.2715987	295	-3.175353	0.0017

Correlation:

	(Intr)	s(*p*T
sin(2 * pi * Time)	0.00	
cos(2 * pi * Time)	-0.06	0.00


Standardized Within-Group Residuals:

Min	Q1	Med	Q3	Max
-2.4500138	-0.6721813	-0.1349236	0.5922957	3.5506618

Number of Observations: 308
Number of Groups: 11

The function ACF allows us to calculate the empirical autocorrelation structure of the residuals from this model:

```
plot(ACF(model),alpha=0.05)
```



You can see that there is highly significant autocorrelation at lags 1 and 2 and marginally significant autocorrelation at lags 3 and 4. We model the autocorrelation structure using one of the standard `corStruct` classes (p. 863). For time series data like this, we typically choose between ‘moving average’, ‘autoregressive’ or ‘autoregressive moving average’ classes. Again, experience with horse biology suggests that a simple moving average model might be appropriate, so we start with this. The class is called `corARMA` and we need to specify the order of the model (the lag of the moving average part). The simplest assumption is that only the first two lags exhibit non-zero correlations ($q=2$):

```
model2 <- update(model, correlation=corARMA(q=2))
anova(model, model2)
```

	Model	df	AIC	BIC	logLik	Test	L.Ratio	p-value
	model	1	5	1669.360	1687.962	-829.6802		
	model2	2	7	1574.895	1600.937	-780.4476	1 vs 2	98.4652 <.0001


This is a great improvement over the original model, which assumed no correlation in the residuals. But what about a different time series assumption? Let us compare the moving average assumption with a simple first-order autoregressive model `corAR1()`:

```
model3 <- update(model2, correlation=corAR1())
anova(model2, model3)
```

	Model	df	AIC	BIC	logLik	Test	L.Ratio	p-value
	model2	1	7	1574.895	1600.937	-780.4476		
	model3	2	6	1562.447	1584.769	-775.2233	1 vs 2	10.4484 0.0012


This is a very significant improvement, $p = 0.0012$, so we choose the `corAR1()` because it has the lowest AIC (it also uses fewer degrees of freedom, d.f. = 6). Error checking on `model3` might proceed like this:

```
plot(model3, resid(., type="p") ~ fitted(.) | Mare)
```



which shows that residuals are reasonably well behaved. And the normality assumption?

```
qqnorm(model3, ~resid(.) | Mare)
```

The errors are close to normally distributed for all of the mares. The model is well behaved, so we accept a first-order autocorrelation structure `corAR1()`.

19.8 Random effects in designed experiments

The rats example, studied by `aov` with an `Error` term on p. 526, can be repeated as a linear mixed-effects model, but only if we recode the factor levels. This example works much better with `lmer` than with `lme`.

```
dd <- read.table("c:\\temp\\rats.txt",h=T)
attach(dd)
names(dd)

[1] "Glycogen" "Treatment" "Rat" "Liver"

Treatment <- factor(Treatment)
Liver <- factor(Liver)
Rat <- factor(Rat)
```

There is a single fixed effect (`Treatment`), and pseudoreplication enters the dataframe because each rat's liver is cut into three pieces and each liver bit is macerated and divided into separate aliquots to produce two readings. First, compute unique factor levels for each rat and each liver bit:

```
rat <- Treatment:Rat
liver <- Treatment:Rat:Liver
```

Then use these as random effects in the `lmer` model:

```
model <- lmer(Glycogen~Treatment+(1|rat)+(1|liver))
summary(model)
```

```
Linear mixed model fit by REML
Formula: Glycogen ~ Treatment + (1 | rat) + (1 | liver)
      AIC      BIC logLik deviance REMLdev
 231.6  241.1 -109.8   234.3    219.6
```

Random effects:

Groups	Name	Variance	Std.Dev.
liver	(Intercept)	14.167	3.7639
rat	(Intercept)	36.065	6.0054
Residual		21.167	4.6007

Number of obs: 36, groups: liver, 18; rat, 6

Fixed effects:

	Estimate	Std. Error	t value
(Intercept)	140.500	4.707	29.850
Treatment2	10.500	6.656	1.577
Treatment3	-5.333	6.656	-0.801

You can see that the treatment effect is correctly interpreted as being non-significant (both t values are less than 2 in absolute value). The variance components (p. 527) can be extracted by expressing the variances as percentages:

```
vars <- c(14.167,36.065,21.167)
100*vars/sum(vars)

[1] 19.84201 50.51191 29.64607
```

Thus 50% of the variation is between rats within treatments, 19.8% is between liver bits within rats and 29.6% is between readings within liver bits within rats. If you are interested principally in the fixed effects, then much the best way to proceed is to average away the pseudoreplication and do a one-way ANOVA with 3 d.f. for error (see p. 525).

19.9 Regression in mixed-effects models

The next example involves a regression of plant size against local point measurements of soil nitrogen (`N`) at five places within each of 24 farms. It is expected that plant size and soil nitrogen will be positively correlated. There is only one measurement of plant size and soil nitrogen at any given point (i.e. there is no temporal pseudoreplication; cf. p. 695):


```
yields <- read.table("c:\\temp\\farms.txt",header=T)
attach(yields)
```

```
names(yields)
```

```
[1] "N" "size" "farm"
```

Here are the data in aggregate, with different plotting colours and symbols for each farm:

```
plot(N,size,pch=rep(16:19,each=40),col=farm)
```



The most obvious pattern is that there is substantial variation in mean values of both soil nitrogen and plant size across the farms: the minimum-yielding fields (yellow) have a mean y value of less than 80, while the maximum (red) fields have a mean y value above 110. Note that because there are more farms (24) than colours (8), we need to use different plotting symbols as well as different colours to distinguish the five points from each farm.

The key distinction to understand is between fitting lots of linear regression models (one for each farm) and fitting one mixed-effects model, taking account of the differences between farms in terms of their contribution to the variance in response as measured by a standard deviation in intercept and a standard deviation in slope. We investigate these differences by contrasting the two fitting functions, `lmList` and `lme`. We begin by fitting 24 separate linear models, one for each farm:

```
linear.models <- lmList(size~N|farm,yields)
coef(linear.models)
```

	(Intercept)	N
1	67.46260	1.5153805
2	118.52443	-0.5550273
3	91.58055	0.5551292
4	87.92259	0.9212662
5	92.12023	0.5380276
6	97.01996	0.3845431
7	68.52117	0.9339957
8	91.54383	0.8220482

```

9      92.04667  0.8842662
10     85.08964  1.4676459
11    114.93449 -0.2689370
12     82.56263  1.0138488
13     78.60940  0.1324811
14     80.97221  0.6551149
15     84.85382  0.9809902
16     87.12280  0.3699154
17     52.31711  1.7555136
18     83.40400  0.8715070
19     88.91675  0.2043755
20     93.08216  0.8567066
21     90.24868  0.7830692
22     78.30970  1.1441291
23     59.88093  0.9536750
24     89.07963  0.1091016

```

You can see very substantial variations in the value of the intercept from 118.52 on farm 2 to 52.32 on farm 17. Slopes are also dramatically different, from negative -0.555 on farm 2 to steep and positive 1.7555 on farm 17. This is a classic problem in regression analysis when (as here) the intercept is a long way from the average value of x (see p. 460); large values of the intercept are almost bound to be correlated with low values of the slope.

Here are the slopes and intercepts from the model specified entirely in terms of random effects: a population of regression slopes predicted within each farm with nitrogen as the continuous explanatory variable, and a population of intercepts for each farm:

```

random.model <- lme(size~1,random=~N|farm)
coef(random.model)

```

```


      (Intercept)          N
1      85.98139    0.574205332
2    104.67366   -0.045401474
3      95.03442    0.331080929
4      98.62679    0.463579847
5      95.00270    0.407906220
6      99.82294    0.207203700
7      85.57345    0.285520355
8      96.09461    0.520896479
9      95.22186    0.672262931
10     93.14157    1.017995748
11    108.27200    0.015213748
12     87.36387    0.689406424
13     80.83933    0.003617022
14     89.84309    0.306402254
15     93.37050    0.636778731
16     92.10914    0.145772156
17     94.93395    0.084935464
18     85.90160    0.709943272
19     92.00628    0.052485987

```

20	95.26296	0.738029408
21	93.35069	0.591151964
22	87.66161	0.673119289
23	70.57827	0.432993929
24	90.29151	0.036747129

Differences between the intercepts explain 97.26% of the variance, differences in slope a mere 0.245%, with a residual variance of 2.49% (see the `summary` table). The thing you notice is that the random effects are less extreme (i.e. closer to the mean) than the fixed effects. This is an example of shrinkage (p. 685), and is clearest from a graphical comparison of the coefficients of the linear and mixed models:

```
mm <- coef(random.model)
ll <- coef(linear.models)
windows(7,4)
par(mfrow=c(1,2))
plot(ll[,1],mm[,1],pch=16,xlab="linear",ylab="random effects")
abline(0,1)
plot(ll[,2],mm[,2],pch=16,xlab="linear",ylab="random effects")
abline(0,1)
```



Most of the random-effects intercepts (left) are greater than their linear model equivalents (they are above the 45 degree line) while most of the random-effects slopes (right) are shallower than their linear model equivalents (i.e. below the line). For farm 17 the linear model had an intercept of 52.317 11 while the random-effects model had an intercept of 94.933 95. Likewise, the linear model for farm 17 had a slope of 1.755 513 6 while the random-effects model had a slope of 0.084 935 465.

We can fit a mixed model with both fixed and random effects. Here is a model in which size is modelled as a function of nitrogen and farm as fixed effects, with farm as a random effect. Because we intend to compare models with different fixed effect structures we need to specify `method="ML"` in place of the default REML.

```
farm <- factor(farm)
mixed.model1 <- lme(size~N*farm,random=~1|farm,method="ML")
mixed.model2 <- lme(size~N+farm,random=~1|farm,method="ML")
mixed.model3 <- lme(size~N,random=~1|farm,method="ML")
mixed.model4 <- lme(size~1,random=~1|farm,method="ML")
anova(mixed.model1,mixed.model2,mixed.model3,mixed.model4)
```

	Model	df	AIC	BIC	logLik	Test	L.Ratio	p-value
mixed.model1	1	50	542.9035	682.2781	-221.4518			
mixed.model2	2	27	524.2971	599.5594	-235.1486	1 vs 2	27.39359	0.2396
mixed.model3	3	4	614.3769	625.5269	-303.1885	2 vs 3	136.07981	<.0001
mixed.model4	4	3	658.0058	666.3683	-326.0029	3 vs 4	45.62892	<.0001

The first model contains a full factorial, with different slopes and intercepts for each of the 25 farms (using up 50 degrees of freedom). The second model has a common slope but different intercepts for the 25 farms (using 27 degrees of freedom); `model2` does not have significantly lower explanatory power than `model1` ($p = 0.2396$). The main effects of farm and of nitrogen application (`model3` and `model4`) are both highly significant ($p < 0.0001$), so we select `model2` because it has the lowest AIC.

Finally, we could do an old-fashioned analysis of covariance, fitting a different two-parameter model to each and every farm without any random effects:

```
model <- lm(size~N*factor(farm))
summary(model)
```

Call:

```
lm(formula = size ~ N * factor(farm))
```

Residuals:

	Min	1Q	Median	3Q	Max
	-3.6077	-1.2947	0.0479	1.0732	4.1297

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	67.46260	14.43749	4.673	1.35e-05	***
N	1.51538	0.73395	2.065	0.0426	*
factor(farm) 2	51.06183	22.86930	2.233	0.0287	*
factor(farm) 3	24.11794	16.54029	1.458	0.1492	
factor(farm) 4	20.45999	34.59610	0.591	0.5561	
factor(farm) 5	24.65762	17.29578	1.426	0.1583	
factor(farm) 6	29.55736	17.74007	1.666	0.1000	
factor(farm) 7	1.05856	20.53771	0.052	0.9590	
factor(farm) 8	24.08122	16.23722	1.483	0.1424	
factor(farm) 9	24.58407	15.45967	1.590	0.1162	
factor(farm) 10	17.62703	16.68467	1.056	0.2943	
factor(farm) 11	47.47189	18.24214	2.602	0.0112	*
factor(farm) 12	15.10002	15.77085	0.957	0.3415	
factor(farm) 13	11.14680	17.82896	0.625	0.5338	
factor(farm) 14	13.50961	19.36739	0.698	0.4877	
factor(farm) 15	17.39122	20.74850	0.838	0.4047	
factor(farm) 16	19.66019	18.72739	1.050	0.2973	
factor(farm) 17	-15.14550	49.01250	-0.309	0.7582	
factor(farm) 18	15.94140	15.15371	1.052	0.2963	
factor(farm) 19	21.45414	17.99214	1.192	0.2370	
factor(farm) 20	25.61956	15.50019	1.653	0.1027	
factor(farm) 21	22.78608	15.65699	1.455	0.1499	
factor(farm) 22	10.84710	17.69820	0.613	0.5419	
factor(farm) 23	-7.58167	16.89435	-0.449	0.6549	

```

factor(farm) 24      21.61703    17.28697     1.250     0.2152
N:factor(farm) 2     -2.07041     0.98369    -2.105     0.0388 *
N:factor(farm) 3     -0.96025     0.89786    -1.069     0.2884
N:factor(farm) 4     -0.59411     1.52204    -0.390     0.6974
N:factor(farm) 5     -0.97735     0.84718    -1.154     0.2525
N:factor(farm) 6     -1.13084     0.97207    -1.163     0.2485
N:factor(farm) 7     -0.58138     0.92164    -0.631     0.5302
N:factor(farm) 8     -0.69333     0.87773    -0.790     0.4322
N:factor(farm) 9     -0.63111     0.81550    -0.774     0.4415
N:factor(farm) 10    -0.04773     0.86512    -0.055     0.9562
N:factor(farm) 11    -1.78432     0.87838    -2.031     0.0459 *
N:factor(farm) 12    -0.50153     0.84820    -0.591     0.5562
N:factor(farm) 13    -1.38290     0.98604    -1.402     0.1651
N:factor(farm) 14    -0.86027     0.89294    -0.963     0.3386
N:factor(farm) 15    -0.53439     0.94640    -0.565     0.5741
N:factor(farm) 16    -1.14547     0.91070    -1.258     0.2125
N:factor(farm) 17     0.24013     1.97779     0.121     0.9037
N:factor(farm) 18    -0.64387     0.79080    -0.814     0.4182
N:factor(farm) 19    -1.31100     0.90886    -1.442     0.1535
N:factor(farm) 20    -0.65867     0.78956    -0.834     0.4069
N:factor(farm) 21    -0.73231     0.81990    -0.893     0.3747
N:factor(farm) 22    -0.37125     0.89597    -0.414     0.6798
N:factor(farm) 23    -0.56171     0.85286    -0.659     0.5122
N:factor(farm) 24    -1.40628     0.95103    -1.479     0.1436

```

```

Residual standard error: 1.978 on 72 degrees of freedom
Multiple R-squared: 0.9678, Adjusted R-squared: 0.9468
F-statistic: 46.07 on 47 and 72 DF, p-value: < 2.2e-16

```

There is a marginally significant overall effect of soil nitrogen on plant size (N has $p = 0.0426$) and (compared to farm 1) farms 2 and 11 have significantly higher intercepts and shallower slopes. The problem, of course, is that this model, with its 24 slopes and 24 intercepts, is vastly overparameterized. Let us fit a greatly simplified model with a common slope but different intercepts for the different farms:

```

model2 <- lm(size~N+factor(farm))
anova(model, model2)

```

Analysis of Variance Table

```

Model 1: size ~ N * factor(farm)
Model 2: size ~ N + factor(farm)
  Res.Df    RSS   Df Sum of Sq    F Pr(>F)
1      72 281.60
2      95 353.81 -23   -72.212 0.8028 0.717

```

This analysis provides no support for any significant differences between slopes. What about differences between farms in their intercepts?

```

model3 <- lm(size~N)
anova(model2, model3)

```

Analysis of Variance Table

```

Model 1: size ~ N + factor(farm)
Model 2: size ~ N
      Res.Df    RSS   Df Sum of Sq      F      Pr(>F)
1         95   353.8         -      -8101.1 94.574 < 2.2e-16 ***
2        118 8454.9  -23      -8101.1 94.574 < 2.2e-16 ***

```

This shows that there are highly significant differences in intercepts between farms. The interpretation of the analysis of covariance is exactly the same as the interpretation of the mixed model in this case where there is balanced structure and equal replication, but `lme` is vastly superior to the linear model when there is unequal replication.

19.10 Generalized linear mixed models

Pseudoreplicated data with non-normal errors lead to a choice of generalized linear mixed-effects models using `lmer` with a specified error family. These were previously handled by the `glmmPQL` function which is part of the `MASS` library (see Venables and Ripley, 2002). That function fitted a generalized linear mixed model with multivariate normal random effects, using penalized quasi-likelihood (hence the ‘PQL’). The default method for a generalized linear model fit with `lmer` has been switched from PQL to the more reliable Laplace method. The `lmer` function can deal with the same error structures as a generalized linear model, namely Poisson (for count data), binomial (for binary data or proportion data) or gamma (for continuous data where the variance increase with the square of the mean). The model call is just like a mixed-effects model but with the addition of the name of the error family, like this:

```
lmer(y~fixed+(time | random), family=binomial)
```

For a worked example with binary data, involving patients who were tested for the presence of a bacterial infection on a number of occasions (the number varying somewhat from patient to patient), see pp. 660–665. The response variable is binary: `yes` for infected patients or `no` for patients not scoring as infected, so the family is binomial. There is a single categorical explanatory variable (a fixed effect) called treatment, which has three levels: drug, drug plus supplement, and placebo. The week numbers in which the repeated assessments on each patient were made is also recorded.

19.10.1 Hierarchically structured count data

This is an example of `lmer` with Poisson errors. Beetles were collected in pitfall traps laid out in a grid of two columns and five rows (10 pitfalls per quadrat) in each of two quadrats, randomly located within each of 3 randomly-located blocks within a field. On each of 4 farms there were 5 protocols of hedgerow management allocated at random to each of fields within the farm (1 = uncut control, 2 = grass cut, 3 = grass cut twice, 4 = hedge cut, 5 = grass and hedge cut). Here are the 1200 counts:

```

data<-read.table("c:\\temp\\nested2.txt",header=T)
attach(data)
head(data)

count farm field block quadrat
1     1     1     1     1       1
2     0     1     1     1       1
3     0     1     1     1       1
4     1     1     1     1       1

```

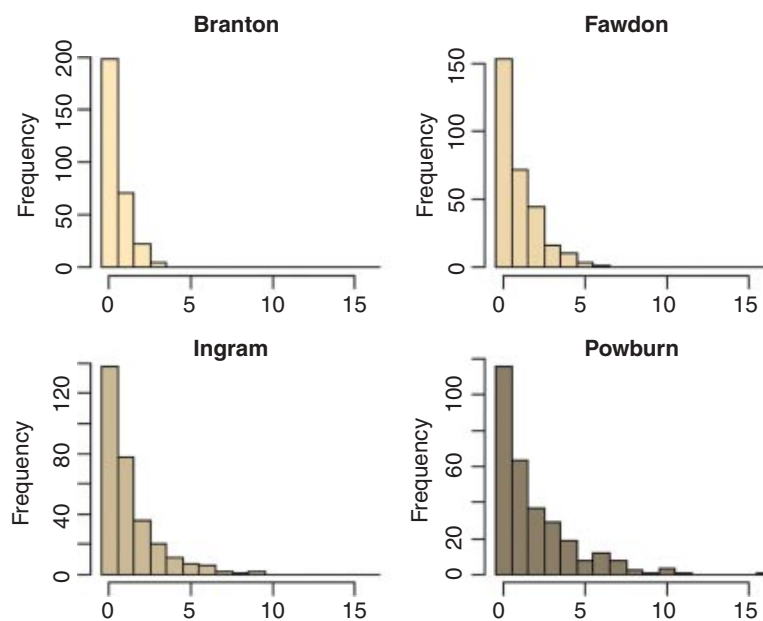


```
5  1  1  1  1  1
6  0  1  1  1  1
```

```
farm<-factor(farm)
field <- factor(field)
block <- factor(block)
quadrat <- factor(quadrat)
library(lme4)
```

This is what the data look like, classified by farms:

```
par(mfrow=c(2,2))
hist(count[farm==1],breaks=-0.5:16.5,
      main="Branton",col="wheat1",xlab="")
hist(count[farm==2],breaks=-0.5:16.5,
      main="Fawdon",col="wheat2",xlab="")
hist(count[farm==3],breaks=-0.5:16.5,
      main="Ingram",col="wheat3",xlab="")
hist(count[farm==4],breaks=-0.5:16.5,
      main="Powburn",col="wheat4",xlab="")
```



Most of the pitfall traps contained no beetles at all sites, and the maximum number caught was 16 on one occasion at Powburn. There are large differences between the farms in mean beetle counts, but despite this, the five hedgerow management treatments showed substantial differences:

```
tapply(count,list(farm,field),mean)

      1      2      3      4      5
1 0.400000 0.466667 0.450000 0.583333 0.583333
2 0.683333 0.766667 0.783333 1.200000 1.133333
3 0.850000 0.916667 1.666667 1.250000 1.383333
4 1.500000 2.216667 1.800000 1.450000 2.200000
```

We need to establish whether the treatment differences (as shown here between the column means) are significant. Because of the massive pseudoreplication, we need to analyse the counts very carefully. We shall use a generalized mixed effects model with Poisson errors, treating field as a fixed effect (the `field` codes refer to the five hedgerow management treatments) with the other factors as nested random effects (quadrats within blocks, within fields, within farms):

```
model <- lmer(count~field+(1|farm/field/block/quadrat),family=poisson)
summary(model)
```

Generalized linear mixed model fit by the Laplace approximation

Formula: count ~ field + (1 | farm/field/block/quadrat)

AIC	BIC	logLik	deviance
2216	2262	-1099	2198

Random effects:

Groups	Name	Variance	Std.Dev.
quadrat:(block:(field:farm))	(Intercept)	0.133785	0.36577
block:(field:farm)	(Intercept)	0.017054	0.13059
field:farm	(Intercept)	0.000000	0.00000
farm	(Intercept)	0.216319	0.46510

Number of obs: 1200, groups: quadrat:(block:(field:farm)), 120;
block:(field:farm), 60; field:farm, 20; farm, 4

Fixed effects:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-0.3322	0.2586	-1.285	0.19889
field2	0.2202	0.1561	1.411	0.15830
field3	0.2851	0.1552	1.836	0.06632 .
field4	0.3205	0.1546	2.074	0.03811 *
field5	0.4402	0.1533	2.872	0.00407 **

As you can see, management treatments 4 and 5 produced significantly higher beetle counts than the controls. Among the random effects you will see that `field:farm` (the fixed effect) registers as zero. Much the biggest cause of variation in beetle numbers was differences between the four farms (variance = 0.216), then differences between the quadrats within a block (0.134), with only very modest differences from block to block within each field (0.017).

If, instead of fitting the management treatment as a fixed effect, we were to fit it as a random effect, then the model changes as follows:

```
model2 <- lmer(count~1+(1|farm/field/block/quadrat),family=poisson)
summary(model2)
```

Generalized linear mixed model fit by the Laplace approximation

Formula: count ~ 1 + (1 | farm/field/block/quadrat)

AIC	BIC	logLik	deviance
2216	2242	-1103	2206

Random effects:

Groups	Name	Variance	Std.Dev.
quadrat:(block:(field:farm))	(Intercept)	0.133591	0.365501
block:(field:farm)	(Intercept)	0.035101	0.187354

```

field:farm          (Intercept) 0.001888 0.043451
farm                (Intercept) 0.213685 0.462260

Number of obs: 1200, groups: quadrat:(block:(field:farm)), 120;
block:(field:farm), 60; field:farm, 20; farm, 4

Fixed effects:
              Estimate Std. Error z value Pr(>|z|)
(Intercept) -0.07779    0.23703  -0.328   0.743

```

The variance component attributable to differences between blocks has almost doubled (this is the spatial scale immediately below fields, which was the scale at which the fixed effects treatments were applied), but the farm-scale variation (larger scale) and quadrat-scale variation (smaller scale) are almost unaffected. So the difference between fields which looks so small as a random effect (0.001 888) is clearly significant when the term is fitted as a fixed effect (as was appropriate in this example).

Here is the much simpler analysis with the pseudoreplication averaged away. Now, there are only 20 numbers in the dataframe (5 numbers (the five treatment mean counts) from each of 4 farms), and we use these to create a new response variable:

```

y <- as.vector(tapply(count,list(farm,field),mean))
y

[1] 0.4000000 0.6833333 0.8500000 1.5000000 0.4666667 0.7666667 0.9166667
[8] 2.2166667 0.4500000 0.7833333 1.6666667 1.8000000 0.5833333 1.2000000
[15] 1.2500000 1.4500000 0.5833333 1.1333333 1.3833333 2.2000000

```

We also need new shorter vectors of categorical explanatory variables for the farm (which now serves as a four-level block in this two-way, non-replicated analysis) and the five-level hedgerow management treatment. The ‘generate factor levels’ function, `gl`, is useful here. For farm blocks (`fblock`) we read the instruction to `gl` like this: ‘generate levels up to 4, each with a repeat of 1, to a total length of 20’. For `hedge` it says ‘generate levels up to 5, each with a repeat of 4’.

```

fblock <- gl(4,1,20)
hedge <- gl(5,4)

```

We cannot fit an exactly analogous model because we no longer have count data (the average beetle numbers are all real numbers) so we cannot use `glm` with Poisson errors. We can get close enough, however, by using a linear model with the logarithms of the mean counts as the response variable like this:

```

model3 <- lm(log(y)~fblock+hedge)
summary(model3)

```

Call:

```
lm(formula = log(y) ~ fblock + hedge)
```

Residuals:

```

      Min       1Q   Median       3Q      Max
-0.29679 -0.08192 -0.01837  0.05106  0.31607

```

Coefficients:

```

              Estimate Std. Error t value Pr(>|t|)
(Intercept)  -0.9555    0.1179  -8.102 3.30e-06 ***
fblock2        0.5943    0.1179   5.039 0.00029 ***
fblock3        0.8728    0.1179   7.400 8.27e-06 ***

```

```
fblock4      1.3008      0.1179  11.030 1.23e-07 ***
hedge2       0.1838      0.1319   1.394 0.18859
hedge3       0.2775      0.1319   2.105 0.05708 .
hedge4       0.3230      0.1319   2.450 0.03060 *
hedge5       0.4383      0.1319   3.324 0.00606 **
```

```
Residual standard error: 0.1865 on 12 degrees of freedom
Multiple R-squared:  0.9214,    Adjusted R-squared:  0.8756
F-statistic:  20.1 on 7 and 12 DF,  p-value: 9.875e-06
```

The p values are slightly different, but the interpretation is exactly the same. After model simplification, it turns out that there is one highly significant contrast, between the last two hedge management treatments (4 and 5) and the rest ($p = 0.0102$). It is hedge cutting, not grass cutting, that makes the difference in beetle counts in this case.

An alternative method of producing the reduced dataframe (i.e. eliminating the pseudoreplication) is to use `aggregate` like this:

```
d2<-aggregate(data,list(farm,field),mean)
model<-lm(log(count)~factor(farm)+factor(field),data=d2)
summary(model)
```

This produces the same output with less than half the R code.